

A prediction algorithm for niche modeling on big environmental data

David R.B. Stockwell*

March 28, 2015

Abstract

This package is an R implementation of the **WhyWhere** data mining algorithm [5]. Developed for biodiversity modelling on big data, the approach is simple and rigorous, efficient to compute, and provides accuracy equal to the best alternative approaches. Thus, it represents the way forward in utilize large data sets of information about species and their environmental relationships.

1 Introduction

This package is concerned with the development of models of species distributions on arbitrarily large environmental data-sets, although the approach could potentially be used in the other fields that require geospatial prediction. Large 2D geospatial data arises from global coverage at fine resolution of temperature, rainfall and other surfaces that might potentially be correlated with some entity of interest.

The operation of the package is similar to the **dismo** package in which implements a number of species distribution models (MaxEnt, Bioclim, Domain, GLM, GAM, and RandomForest). It also follows the pattern of such sampling background points, and the program flow as developed in [4] and commonly used for species distribution modeling.

The novel contribution of WhyWhere is to implement an algorithm that addressed the big data problem, as the size of data sets expand there is a need to reexamine traditional approaches to data analysis with expensive memory or computational demands and if necessary redesign the algorithms that in order to achieve the same purposes. A number of features were proposed in the WhyWhere algorithm in [5] to remedy issues that often block or limit application of other algorithms in some way:

*All correspondence to be addressed to the author at david.r.stockwell@cqu.edu.au, Adjunct Researcher at Central Queensland University.

- 31 1. Memory limitations: No more than one environmental data set in memory
32 at one time. In most the **dismo** package all methods take prepared data
33 in a 'wide' file format, i.e. similar to an excel spreadsheet. For this the
34 data is extracted from a raster data brick. All environmental data must
35 be in memory for the data brick which can exceed the memory limitations
36 of the machine.
- 37 2. Incompatible resolutions in environmental data sets: In the **raster** package
38 data coarse scale data-sets are enlarged to have the same resolution in
39 order to form a data "brick" of coordinated grids of the same resolution.
40 This is approach is memory expensive and is avoided in **WhyWhere**.
- 41 3. Mixed types and distributions: Robust to distributional properties. Some
42 methods such as generalized linear regression (GLMs) are very sensitive
43 to distributional assumptions. The method in **WhyWhere** can also handle
44 factor variables. Few algorithms can handle both continuous and categor-
45 ical variables. **WhyWhere** is similar to MaxEnt in this respect.

46 A number of advances have been implemented since the original **WhyWhere**
47 implementation. We make use of the **raster** package for handling gridded spatial
48 data [6] and also the **data.table** package for fast aggregation of large data [1].
49 The model is developed to optimize the Area Under Curve statistic rather than
50 a Chi-squared statistic. In the original WhyWhere [5] a pcolor quantization al-
51 gorithm was employed to segment the environmental space into groups in order
52 to assign predicted probabilities for each group. Higher dimension models are
53 now developed by combining vectors using fuzzy set conjunction (minimum) or
54 disjunction (maximum). The approach is quite natural to predicting distribu-
55 tions on large numbers of large files, and the modifications are even more similar
56 to machine learning approach.

57 The guide is arranged as follows. Section 2 illustrates the basic usage. Sec-
58 tion 3 describes the theoretical basis. Section 4 compares the performance with
59 the popular MaxEnt method on well known data-sets.

60 2 Example

61 The species distribution models are typically built from two data sources: the
62 coordinates of the species of interest, and the environmental values, obtained as
63 2D map layers. The function `whywhere` is a wrapper around a number of more
64 basic functions to perform the entire analysis in one step. Below illustrates a
65 complete analysis on the *Bradypus* file of the **dismo** package [3] with output in
66 Figure 1.

67 The inputs to **WhyWhere** are as follows: a **data.table** with the longitude
68 and latitude of known locations, and a list of environmental data files that may
69 be read into the raster package. A couple of other parameters are available:
70 *multi* for conducting the search in multiple dimensions, *limit* for the minimum
71 for entry into the results, *beam* to specify the number of entries to keep in the

```

> source("../R/ww2.R")
> library(dismo)
> #data(wrld_simpl)
> files <- list.files(path=paste(system.file(package="dismo"),'/ex',sep=''), pattern='grd',
> file <- paste(system.file(package="dismo"), '/ex/bradypus.csv',sep='')
> Pres <- read.table(file, header=T,sep=",")
> Pres$species=NULL
> result=whywhere(Pres,files)
> plot.model(result)

```




Figure 1: **WhyWhere** predicted distribution of the Bradypus data set.

72 beam search table, e the size of border around presence points, and *plot* if you
 73 want to plot as you go. The result is a list of components: results, lookup, etc.
 74 The main algorithm implements a classic beam search as follows:

- 75 1. label input coordinates with 1 and find range
- 76 2. generate random points within this range and label with 0
- 77 3. for all environmental files do
 - 78 (a) develop membership function for file
 - 79 (b) bind variable name and AUC to result
 - 80 (c) test conjunct of this variable with best
 - 81 (d) if result is better then bind to list
 - 82 (e) delete last row from result if too many
- 83 4. output

84 The two analytic steps of interest are the membership function and the
 85 combination of variables into expressions. The membership function takes the
 86 following inputs and outputs model data: *file* is a single environmental file, *ext*
 87 is the geographic extent, presence and absence locations, *pa* a vector of zeros and
 88 ones for presence and absence locations, and returns a model object *membership*.
 89 The model consists of a table of ranges of the variable, or factors for a categorical
 90 variable, with the counts and posterior probability in each category.

91 Figure 2 shows a table output for the highest rated variable in the Bradypus
 92 dataset. The categorical ranges are listed on the x axis, the prior distribution
 93 (background or 0s) is the black solid line, and the distribution of presences (1s)
 94 in each category is the dashed line. The posterior probability is the product
 95 of these (gray bars). Note that the almost uniform distribution of background
 96 classes due to the quantile cuts.

97 The procedure in membership is as follows:

- 98 1. get the file in raster format
- 99 2. crop according to extent
- 100 3. extract environmental values from the file
- 101 4. determine number of breaks
- 102 5. cut into variable into quantiles (or factors)
- 103 6. construct table of factors with number of counts in each
- 104 7. calculate posterior probability
- 105 8. label row with probability

```
> plot.membership(result)
```

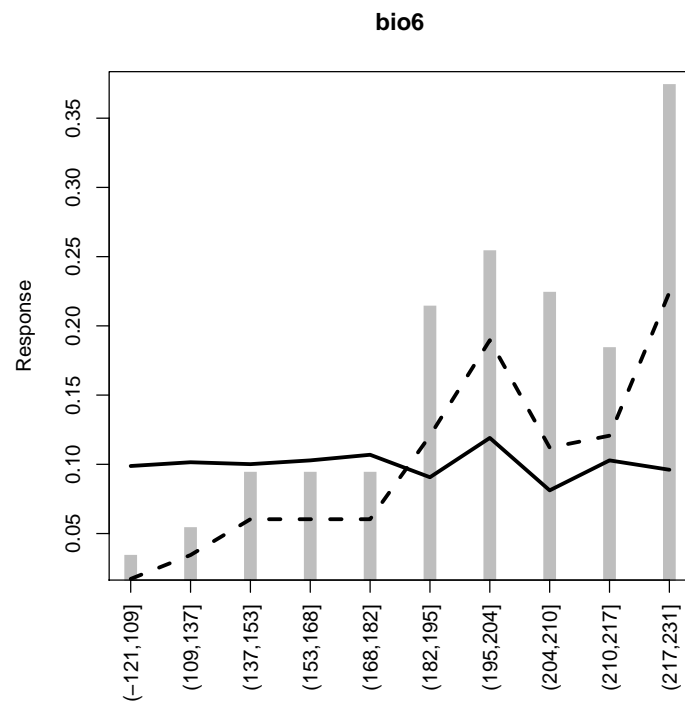


Figure 2: Response function bio6 plotted.

- 106 9. calculate AUC
- 107 10. return vector of items labeled with probability

108 Returning to the main algorithm, calculated before calling membership: the
 109 geographic extent, and the random sample of background data. This is necessary
 110 as the clip and background points must not vary. If the supplied species data
 111 does not contain background data, absences, then they are generated.

112 The membership function is used to produce a vector of probabilities. Vec-
 113 tors of probabilities for each variable are then combined using the fuzzy maxi-
 114 mum principle to maximize the AUC. The rationale for this approach and com-
 115 parison is in the following section. The output is a table of the best d variables
 116 as indicated by the AUC.

117 This scheme allows only one environmental variable to be loaded at a time.
 118 The rasters may have different resolutions, as the environmental values are ex-
 119 tracted only in the one environmental variable. The variable is also cropped
 120 at this time. We also test a combination of the variables with the next best
 121 variables by applying the minimum of the item probability vectors and recalcu-
 122 lating the AUC. It is possible to monitor the progress of **WhyWhere** with the
 123 plot option. This plots out the best model sofar and prints out a list of the best
 124 models considered.

125 3 Theory

126 Elements of justification of the approach are spread throughout the literature.
 127 The basic approach, that of segmentation of the continuous variables into equal
 128 portions has been around since early image processing days. The GIF format
 129 would compress images by reducing the number of colors in the red, green,
 130 blue 3D space. Tests found superiority of the median cut algorithm, for color
 131 reduction where each class has an equal number [2] . Here we quantiles for the
 132 same reason.

133 Median cuts have a good statistical basis as uniform sizes minimize the
 134 variability of the estimates of the posterior probability. Experiments performed
 135 with other schemes such as equal cuts were inferior. The exception is the factor
 136 variable which uses only those values in the layer. When using factor variables
 137 all factors are labeled as such.

138 The benefit of the approach is that it does not make assumptions about
 139 the distribution. Many make inappropriate assumptions about the distribution,
 140 even linear or monotonically increasing. In ecological niche theory a species is
 141 more abundant, or be more likely, within a limited range of an environmental
 142 variable. The typical variables used are temperature and precipitation ranges,
 143 and here a uni-modal or humped distribution is the simplest viable distribution.
 144 In addition, many environmental variables are categorical variables such as soil
 145 and vegetation type, so the approach is capable of handling the true relationships
 146 of species to environment without making assumptions about them.

147 The use of the quotient of counts needs justification. The locations where
148 the species occurs can be thought of as a sampling of the environmental space.
149 That is, while the counts of the environmental values has a distribution P_0 for
150 over the range of the variable, the distribution of the sample where the species is
151 present is P_1 over the same values. The most significant variable has the greatest
152 difference between these distributions. The membership function is developed
153 from the ration of counts P_1/P_0 .

154 When we look at the change in this distribution between the background
155 counts and the counts of the species sample, the classic approach to evaluating
156 the significance of the the Chi-squared test, and this is what was used in the
157 original **WhyWhere** package. The K-S test can also be used. However Chi2
158 doesn't work in evaluating the multivariate case.

159 Turns out we don't need it. If we work out the probability for each of
160 the training data points and apply fuzzy set methodology to these vectors the
161 resulting vector of probabilities can be evaluated using the AUC.

162 In another derivation the membership function can be viewed as a relaxation
163 of some aspects of strict Bayesian statistics.

$$P(S|E) = P(E|S)P(S)/P(E) \quad (1)$$

164 But the problem is ratios of counts are not strictly probabilities. In any
165 sense, probability of the occurrence of a species $P(S)$ is not well defined, or not
166 generally of interest. In a typical example the probably of finding a species is
167 dependent on season, search effort, and so is not well controlled. Usually we
168 only want to know the best areas to search for a species, as predicted by the
169 most significant environmental variables. So dropping the $P(S)$ we still have a
170 proportional relationship which sufficient to compare alternative environmental
171 variables.

$$P(S|E) \propto P(E|S)/P(E) \quad (2)$$

172 Other similar developments are using in MaxEnt where the membership is
173 done using more complicated. The results are very similar as the final compar-
174 ative section shows.

175 3.1 Fuzzy Conjunctions

176 The inferential basis is fuzzy set theory, where instead statements that are either
177 true or false, a membership function describes a fuzzy truth value as a function
178 $f : \mathbb{R} \rightarrow [0, 1]$ from a variable V to the real unit interval $[0, 1]$. One must consider
179 membership functions taking values from other spaces such as categories, (also
180 known as factors in R) N or on a space of many variables $V_1 \times V_2 \times \dots \times V_n$
181 where each V_i is an interval in N or R .

182 Experience has shown that a particularly useful way of combining unitary
183 membership functions to resemble the AND, OR and NOT operators of classical
184 logic are Zadeh operators:

185 AND: $x \wedge y = \min(f(x), f(y))$,
186 OR: $x \vee y = \max(f(x), f(y))$, and
187 NOT: $\neg x = (1 - f(x))$.

188 There have been many approaches to learning fuzzy rules from from given
189 data, and approaches to representation of the discovered rules. One of the
190 outstanding problems is the trade-off between accuracy and interpretability, or
191 prediction and explanation in the ecological literature. In particular in this pack-
192 age, ecological theory can motivate our approach, thus satisfying both criteria.
193 It is the desire to address the problem of explanatory models that motivated t
194 development of **WhyWhere** – to describe Why is a species Where? – without
195 sacrificing predictive accuracy or computational tractability.

196 In the one dimensional case the variable with the largest significance is the
197 best to select. This gives a list of variables that the species responds to the
198 strongest.

199 But when we add more responses into a fuzzy expressions in a multi-variable
200 data-sets the most accurate conjunctive expression may not contain the best
201 single variable. This means we cannot use such methods as greedy search to
202 identify higher dimensional expressions. The problem of expressing environ-
203 mental relationships of more than one variable has been an important topic in
204 statistical and ecological research, generalized linear modelling machine learn-
205 ing on the other. Here we need to consider the logical relationship between
206 variables that are expected from ecological theory.

207 3.1.1 Law of the Minimum

208 The principle of Liebig’s Law of the Minimum states that growth is controlled
209 by the scarcest necessary resource. This is logically an AND operation or a fuzzy
210 conjunction of limiting factors. Note that a model composed of an arithmetic
211 sum would represent the concept of growth determined by the overall sum of
212 resources contributed from different sources, and so is not consistent with the
213 Law of the Minimum. Most model used to predict species distributions are of
214 this form, and so their capacity to explain species distributions is questionable.

215 3.1.2 Competitive exclusion principle

216 In ecology, Gauss’s law of competitive exclusion is a proposition that two species
217 competing for the same resource cannot coexist at constant populations if all
218 other things remain equal, due to effect of slight advantages magnified over
219 generations. This behavioral shift leads to ecological niches. This is logically an
220 OR operation or a fuzzy disjunction:

221 A possible example of this case is where a set of location actually contains two
222 different species with different different habitats. A disjunction of two habitats
223 may model this case better.

224 To evaluate both unit variate and multivariate combination of environmental
225 variables, we predict the probability of presence on each location and calculate
226 the area under the curve of the receiver operating statistic (AUC for short).

227 There is a further statistic that is the area under rank correlation or AUC.
228 This arises where a models outputs a figure in the range of zero to one, but
229 a prediction must be made of 0 or 1. We must select a cutoff value for the
230 prediction to assign to zero or one. The AUC is the probability that would use
231 the optimal accuracy.

232 A high AUC indicates that sites with high membership are more likely to
233 be areas of presence, and vice versa. An AUC score of 0.5 is no better than
234 random. The AUC can be calculated from.

235 4 Comparison

236 We test other models against the response output for the single dimensional
237 case. The response functions for MaxEnt and **WhyWhere** are quite similar as
238 follows - when we can get the rgdal package installed!!!

239 5 Conclusion and Further Work

240 The R package **WhyWhere** is a useful packages to fit, plot and test empirical
241 species as a conjunction of response functions. More complex logical expressions
242 are planned, as are improvements in computational efficiency and access to
243 cloud data. The method provides simple and intuitive, efficient to compute,
244 and typical predictive results that are at least equal to the best alternative
245 approaches.

246 References

- 247 [1] M Dowle, T Short, S Lianoglou, and A Srinivasan. data.table: Extension of
248 data.frame. 10 2014.
- 249 [2] Paul Heckbert. Color image quantization for frame buffer display. In *SIG-*
250 *GRAPH '82: Proceedings of the 9th annual conference on Computer graphics*
251 *and interactive techniques*, pages 297–307, New York, NY, USA, 1982. ACM
252 Press.
- 253 [3] R.J. Hijmans, S. Phillips, J. Leathwick, and J. Elith. dismo: Species distri-
254 bution modeling with r. 2011.
- 255 [4] D Stockwell and D Peters. The GARP modelling system: problems and solu-
256 tions to automated spatial prediction. *International Journal Of Geographical*
257 *Information Science*, 13(2):143–158, 1999.
- 258 [5] David R B Stockwell. Improving ecological niche models by data mining
259 large environmental datasets for surrogate models. *Ecological Modelling*,
260 192:188–196, 2006.

- 261 [6] Robert J. Hijmans & Jacob van Etten. raster: Geographic analysis and
262 modeling with raster data. 2012. R package version 2.0-12.